

# Project\_Final

October 3, 2017

## 1 School Finances Dataset

```
In [1]: import missingno as msno
import seaborn as sns
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
sns.set()
import math
import csv
from scipy import stats
```

```
from sklearn.preprocessing import MinMaxScaler
import statsmodels.api as sm
import matplotlib.mlab as mlab
from sklearn.feature_selection import RFECV
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
```

```
/usr/lib/python3.4/importlib/_bootstrap.py:321: FutureWarning: The pandas.core.datetools module
return f(*args, **kwds)
```

```
In [2]: df = pd.read_csv('elsec15.csv')
```

```
/usr/local/lib/python3.4/dist-packages/IPython/core/interactiveshell.py:2698: DtypeWarning: Colu
interactivity=interactivity, compiler=compiler, result=result)
```

### 1.1 Understanding the Data

```
In [3]: df.shape
```

```
Out[3]: (14376, 141)
```

```
In [4]: df.head()
```

```

Out [4]:      STATE      IDCENSUS      NAME  CONUM  CSA   CBSA  \
0         1  1500100100000  AUTAUGA COUNTY SCHOOL DISTRICT  1001    N  33860
1         1  1500200100000  BALDWIN COUNTY SCHOOL DISTRICT  1003  380  19300
2         1  1500300100000  BARBOUR COUNTY SCHOOL DISTRICT  1005    N      N
3         1  1500300200000    EUFAULA CITY SCHOOL DISTRICT  1005    N      N
4         1  1500400100000    BIBB COUNTY SCHOOL DISTRICT  1007  142  13820

      SCHLEV  NCESID  YRDATA   V33  ...   V32   _19H   _21F   _31F   _41F  \
0         3  0100240    15  9664  ...    0  49431  16603  2992  63042
1         3  0100270    15  30596  ...    0  337160  99087  13027  423220
2         3  0100300    15   925  ...    0   8024    0    304   7720
3         3  0101410    15  2829  ...    0     0    0     0     0
4         3  0100360    15  3357  ...    0  22155    0   1190  20965

      _61V  _66V   W01   W31   W61
0         0    0  2094   372  8617
1         0    0  5784  50441  71370
2         0    0    0     0   646
3         0    0    0  2054  7478
4         0    0  1397   790  5400

```

[5 rows x 141 columns]

```
In [5]: df.describe()
```

```

Out [5]:      STATE      IDCENSUS      CONUM      SCHLEV  YRDATA  \
count  14376.000000  1.437600e+04  14376.000000  14376.000000  14376.0
mean     26.801336  2.728090e+13  29838.158598     2.883139    15.0
std     13.894331  1.389514e+13  14753.492121     1.271649     0.0
min       1.000000  1.500100e+12   1001.000000     1.000000    15.0
25%      15.000000  1.550327e+13  18063.000000     3.000000    15.0
50%      27.000000  2.750320e+13  30063.000000     3.000000    15.0
75%      38.000000  3.850053e+13  41009.000000     3.000000    15.0
max       51.000000  5.150230e+13  56045.000000     7.000000    15.0

      V33      TOTALREV      TFEDREV      C14  \
count  14376.000000  1.437600e+04  1.437600e+04  14376.000000
mean     3374.682526  4.546951e+04  3.708038e+03   915.914858
std    14419.737037  2.590623e+05  2.021328e+04  5972.748111
min         0.000000  0.000000e+00  0.000000e+00    0.000000
25%       305.000000  5.241500e+03  2.920000e+02   54.000000
50%       979.500000  1.402300e+04  8.470000e+02  180.000000
75%      2744.250000  3.791375e+04  2.433500e+03  545.250000
max    995192.000000  2.543738e+07  1.307783e+06  379531.000000

      C15      ...      V32      _19H      _21F  \
count  14376.000000  ...  14376.000000  1.437600e+04  1.437600e+04
mean     752.781093  ...      7.318656  2.871850e+04  4.831726e+03

```

std	3619.184513	...	122.649299	1.841763e+05	2.652979e+04
min	0.000000	...	0.000000	0.000000e+00	0.000000e+00
25%	0.000000	...	0.000000	8.400000e+01	0.000000e+00
50%	118.000000	...	0.000000	4.034000e+03	0.000000e+00
75%	532.000000	...	0.000000	1.866350e+04	3.030000e+02
max	248209.000000	...	7753.000000	1.372802e+07	1.312286e+06

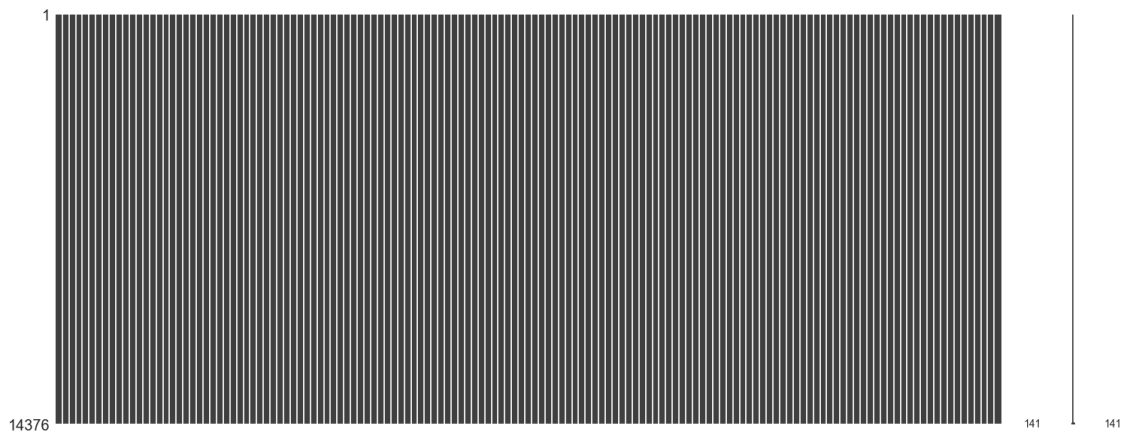
	_31F	_41F	_61V	_66V \
count	14376.000000	1.437600e+04	14376.000000	14376.000000
mean	3936.853158	2.957606e+04	487.162354	553.761547
std	20071.591744	1.894224e+05	3555.988118	6968.740454
min	0.000000	0.000000e+00	0.000000	0.000000
25%	5.000000	7.150000e+01	0.000000	0.000000
50%	410.000000	4.143500e+03	0.000000	0.000000
75%	1845.000000	1.912600e+04	0.000000	0.000000
max	731854.000000	1.447108e+07	173300.000000	700000.000000

	W01	W31	W61
count	14376.000000	14376.000000	1.437600e+04
mean	1335.636825	3710.651920	9.091542e+03
std	10651.097318	20183.852905	3.261580e+04
min	0.000000	0.000000	0.000000e+00
25%	0.000000	0.000000	8.350000e+02
50%	0.000000	0.000000	2.708500e+03
75%	491.000000	626.250000	7.567500e+03
max	869643.000000	885058.000000	2.355662e+06

[8 rows x 137 columns]

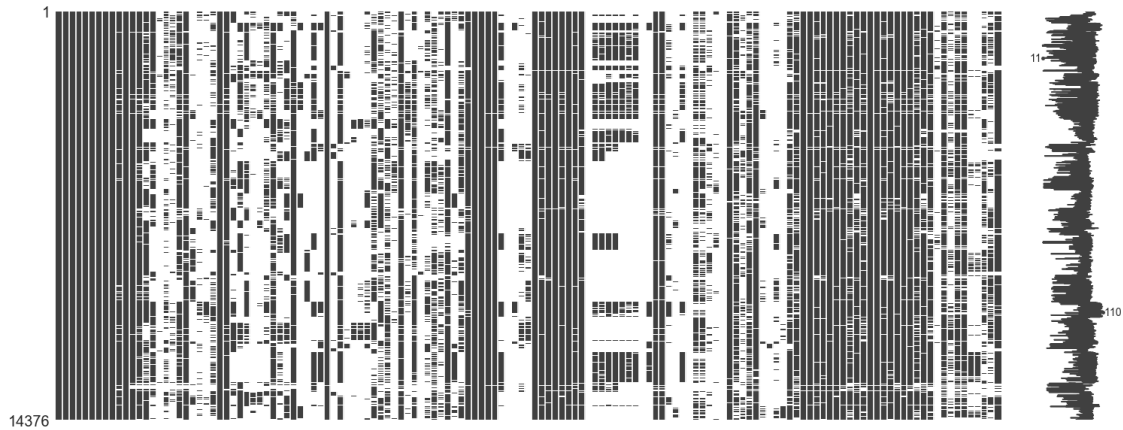
What does our Matrix look like?

In [6]: `msno.matrix(df)`



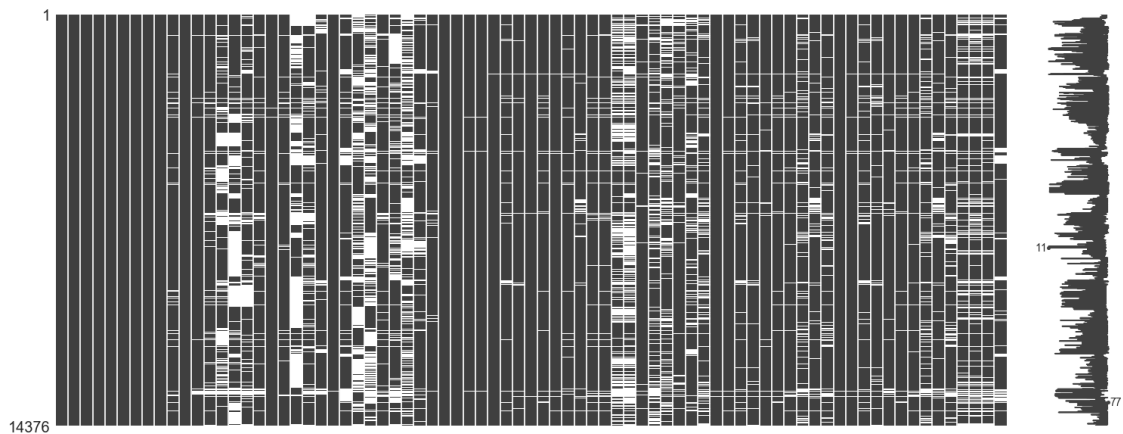
So all the data is here. However some columns have a lot of 0's. Lets see this

```
In [7]: df1 = df.replace(0, float("NaN"))
        msno.matrix(df1)
```



So alot of the data is 0's. Lets see what the data looks like after removing columns have > 50% of their entries as 0's.

```
In [8]: df1 = df1.dropna(axis=1,thresh=math.ceil(df.shape[0]/2), how='all')
        msno.matrix(df1)
        print("%s columns have more than 50%% of their entries as 0's." % str(df.shape[1] - df1.shape[1]))
```



64 columns have more than 50% of their entries as 0's.

So about half the columns have > 50% of their entries as zeroes.  
How many rows have > 50% of their entries as zeroes?

```
In [9]: df1 = df.replace(0, float("NaN"))
        df1 = df1.replace("N", float("NaN"))
        df1 = df1.dropna(axis=0,thresh=math.ceil(df1.shape[1]/2), how='all')
        print("%s rows have more than 50%% of their entries as 0's." % str(df.shape[0] - df1.shape[0]))
```

2432 rows have more than 50% of their entries as 0's.

```
In [10]: df.columns
```

```
Out[10]: Index(['STATE', 'IDCENSUS', 'NAME', 'CONUM', 'CSA', 'CBSA', 'SCHLEV', 'NCESID',  
               'YRDATA', 'V33',  
               ...  
               'V32', '_19H', '_21F', '_31F', '_41F', '_61V', '_66V', 'W01', 'W31',  
               'W61'],  
              dtype='object', length=141)
```

## 1.2 Cleaning the Data

We have decided to work only with the columns that have > 50 % of the rows as non-zero entries. We thought that if most of the columns are zero, there is not enough discrepancy in the data to make an interesting analysis.

```
In [11]: df = df.replace(0, float("NaN"))  
         df = df.dropna(axis=1, thresh=math.ceil(df.shape[0]/2), how='all')  
         df = df.replace(float("NaN"), 0)
```

Additionally some of the data is not very useful to us. For example the census ID's. We have omitted these columns also.

```
In [12]: df.drop("YRDATA", axis=1, inplace=True)      #Its all 2015  
         df.drop("IDCENSUS", axis=1, inplace=True)  
         df.drop("NCESID", axis=1, inplace=True)
```

No data is missing, so no data needs to be imputed. We made some mappings between column values and human understandable values, for understanding more of the data.

```
In [13]: State_Dict = {1: 'AL', 2: 'AK', 3: 'AZ', 4: 'AR', 5: 'CA', 6: 'CO', 7: 'CT', 8: 'DE', 9:  
  
         Col_To_English = {}  
         #This is a local file referenced.  
         with open('School Data Labels.csv', 'r') as csvfile:  
             csvreader = csv.reader(csvfile, delimiter=',')  
             for row in csvreader:  
                 Col_To_English[row[0]] = row[1]  
  
         for key in Col_To_English:  
             print("Code: ", key, "| Label: ", Col_To_English[key])
```

```
Code: V29 | Label: Total salaries and wages - Food services  
Code: D11 | Label: Revenue from other school systems  
Code: Z34 | Label: Total employee benefit payments  
Code: TOTALREV | Label: TOTAL ELEMENTARY-SECONDARY REVENUE  
Code: V10 | Label: Total employee benefit payments - Instruction
```

Code: Z33 | Label: Total salaries and wages - Instruction  
 Code: TOTALEXP | Label: TOTAL ELEMENTARY-SECONDARY EXPENDITURE  
 Code: U97 | Label: Miscellaneous other local revenues  
 Code: Z32 | Label: Total salaries and wages  
 Code: C13 | Label: All other revenues from state sources  
 Code: E13 | Label: Current operation expenditure - Instruction  
 Code: V13 | Label: Total salaries and wages - Instructional staff support  
 Code: TFEDREV | Label: TOTAL ELEMENTARY-SECONDARY REVENUE  
 Code: TCUROTH | Label: TOTAL CURRENT SPENDING FOR OTHER ELEMENTARY-SECONDARY PROGRAMS  
 Code: K10 | Label: Other equipment  
 Code: C15 | Label: Federal revenue through the state - Children with disabilities - IDEA  
 Code: TCURINST | Label: TOTAL CURRENT SPENDING FOR INSTRUCTION  
 Code: SCHLEV | Label: School Level Code  
 Code: TCAPOUT | Label: TOTAL CAPITAL OUTLAY EXPENDITURE  
 Code: E08 | Label: Current operation expenditure - General administration  
 Code: K09 | Label: Instructional equipment  
 Code: NAME | Label: School System Name  
 Code: TCURELSC | Label: TOTAL CURRENT SPENDING FOR ELEMENTARY-SECONDARY PROGRAMS  
 Code: V16 | Label: Total employee benefit payments - General administration  
 Code: V90 | Label: Current operation expenditure - Business/central/other support services  
 Code: V38 | Label: Total employee benefit payments - Business/central/other support services  
 Code: E11 | Label: Current operation expenditure - Food services  
 Code: V23 | Label: Total salaries and wages - Student transportation  
 Code: C10 | Label: School lunch programs  
 Code: F12 | Label: Construction  
 Code: V11 | Label: Total salaries and wages - Pupil support  
 Code: I86 | Label: Interest on school system debt  
 Code: W61 | Label: Cash and deposits, held at end of fiscal year - Other funds  
 Code: E07 | Label: Current operation expenditure - Instructional staff support  
 Code: V17 | Label: Total salaries and wages - School administration  
 Code: V22 | Label: Total employee benefit payments - Operation and maintenance of plant  
 Code: V15 | Label: Total salaries and wages - General administration  
 Code: A09 | Label: School lunch revenues  
 Code: TSTREV | Label: Total Revenue from State Sources  
 Code: C01 | Label: General formula assistance  
 Code: V18 | Label: Total employee benefit payments - School administration  
 Code: \_31F | Label: Long-term debt retired during the fiscal year  
 Code: C20 | Label: Federal revenue through the state - All other  
 Code: Q11 | Label: Payments to other school systems  
 Code: \_19H | Label: Long-term debt outstanding at beginning of the fiscal year  
 Code: V40 | Label: Current operation expenditure - Operation and maintenance of plant  
 Code: C16 | Label: Federal revenue through the state - Math, science, and teacher quality  
 Code: TLOCREV | Label: Total Revenue from Local Sources  
 Code: NONELSEC | Label: TOTAL CURRENT SPENDING FOR NONELEMENTARY-SECONDARY PROGRAMS  
 Code: \_41F | Label: Long-term debt outstanding at end of fiscal year  
 Code: V37 | Label: Total salaries and wages - Business/central/other support services  
 Code: CONUM | Label: ANSI State and County Code  
 Code: V45 | Label: Current operation expenditure - Student transportation

```

Code: STATE | Label: State Identification Number
Code: T06 | Label: Property taxes
Code: V24 | Label: Total employee benefit payments - Student transportation
Code: TCURSSVC | Label: TOTAL CURRENT SPENDING FOR SUPPORT SERVICES
Code: V33 | Label: Fall Membership
Code: CBSA | Label: Core-Based Statistical Area
Code: V14 | Label: Total employee benefit payments - Instructional staff
Code: E09 | Label: Current operation expenditure - School administration
Code: V30 | Label: Total employee benefit payments - Food services
Code: V21 | Label: Total salaries and wages - Operation and maintenance of plant
Code: C14 | Label: Federal revenue through the state - Title I
Code: U22 | Label: Interest earnings
Code: V12 | Label: Total employee benefit payments - Pupil support
Code: A13 | Label: District activity receipts

```

### 1.3 Visualizations

First lets look at the geographic distribution of our data

```

In [14]: State_Dict = {1: 'AL', 2: 'AK', 3: 'AZ', 4: 'AR', 5: 'CA', 6: 'CO', 7: 'CT', 8: 'DE', 9:
State_Rows = []
for key in State_Dict:
    mask = (df["STATE"] == key)
    States_Values = df[mask]
    count = mask.sum()
    # Corruption is a data value dealt with later
    corruption = (States_Values["V15"].sum() + States_Values["V16"].sum() + States_Valu

    mydict = {"State": State_Dict[key], "Count": count, "Corruption":corruption}
    State_Rows.append(mydict)
States = pd.DataFrame(State_Rows)

In [15]: # GENERATES OUR NUMBER OF SCHOOLS PER STATE CHLOROPLETH
import plotly.plotly as py # Note you need to setup an API key for this to work.
from plotly.graph_objs import *

scl = [[0.0, 'rgb(252,240,247)'],[0.2, 'rgb(218,218,235)'],[0.4, 'rgb(188,189,220)'],\
[0.6, 'rgb(158,154,200)'],[0.8, 'rgb(117,107,177)'],[1.0, 'rgb(84,39,143)']]
data = [ dict(
    type='choropleth',
    colorscale = scl,
    autocolorscale = False,
    locations = States['State'],
    z = States['Count'].astype(float),
    locationmode = 'USA-states',
    marker = dict(
        line = dict (

```

```

        color = 'rgb(255,255,255)',
        width = 2
    ) ),
    colorbar = dict(
        title = "Number of Schools")
    ]
layout = Layout(
    autosize=False,
    geo=dict(
        countrycolor='rgb(102, 102, 102)',
        countrywidth=0.1,
        lakecolor='rgb(255, 255, 255)',
        landcolor='rgba(237, 247, 138, 0.28)',
        lonaxis=dict(
            gridwidth=1.5999999999999999,
            range=[-180, -50],
            showgrid=False
        ),
        projection=dict(
            type='albers usa'
        ),
        scope='usa',
        showland=True,
        showrivers=False,
        showsubunits=True,
        subunitcolor='rgb(102, 102, 102)',
        subunitwidth=0.5
    ),
    hovermode='closest',
    images=list([
        dict(
            x=1,
            y=0.6,
            sizex=0.155,
            sizey=0.4,
            source='http://i.imgur.com/Xe3f1zg.png',
            xanchor='right',
            xref='paper',
            yanchor='bottom',
            yref='paper'
        )
    ]),
    showlegend=True,
    title='<b>Number of Schools per State</b>',
    width= 800,
    margin = dict(
        l=0,
        r=50,

```



```

        b=100,
        t=100,
        pad=4)
    )
    fig = Figure(data=data, layout=layout)
    py.ipplot(fig, filename='DSD_1')

```

Out[15]: <plotly.tools.PlotlyDisplay object>

Now we are going to add a column which converts enrollment size to a categorical data type. Note that the average enrollment size is 693, according to <https://nces.ed.gov/pubs2012/2012001.pdf>, Page 170.

```

In [16]: def group_pop(x):
        if x >= 50000:
            return "Greater than 50,000"
        elif x < 50000 and x >= 10000:
            return "Greater than 9,999 and less than 50,000"
        elif x < 10000 and x >= 2000:
            return "Greater than 1999 and less than 10,000"
        elif x < 2000 and x >= 999:
            return "Greater than 999 and less than 2000"
        elif x < 1000 and x >= 499:
            return "Greater than 499 and less than 1000"
        elif x < 500 and x >= 99:
            return "Greater than 99 and less than 500"
        else:
            return "Less than 100"

df["Enrollment Size Class"] = df['V33'].map(group_pop)
df_pop_mean = df.groupby(df["Enrollment Size Class"]).count()
df_pop_sum = df.groupby(df["Enrollment Size Class"]).sum()
df_pop_count= df.sort_values("V33", ascending=True).groupby(df["Enrollment Size Class"])

In [17]: plt.figure(figsize=(11,11))

```

```

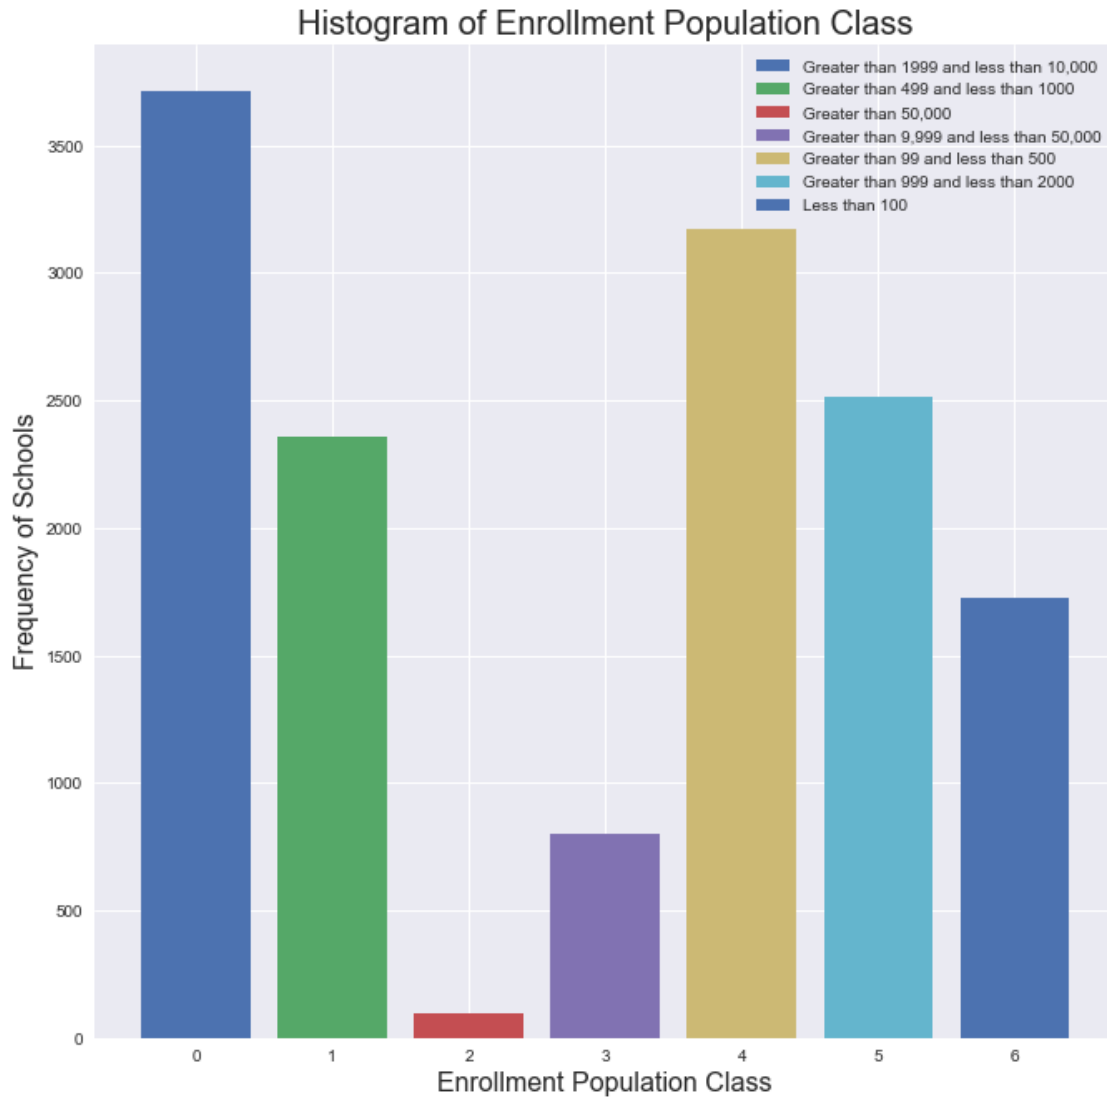
plt.xlabel('Enrollment Population Class', fontsize=16)
plt.ylabel('Frequency of Schools', fontsize=16)
plt.title('Histogram of Enrollment Population Class', fontsize=20)

```

```

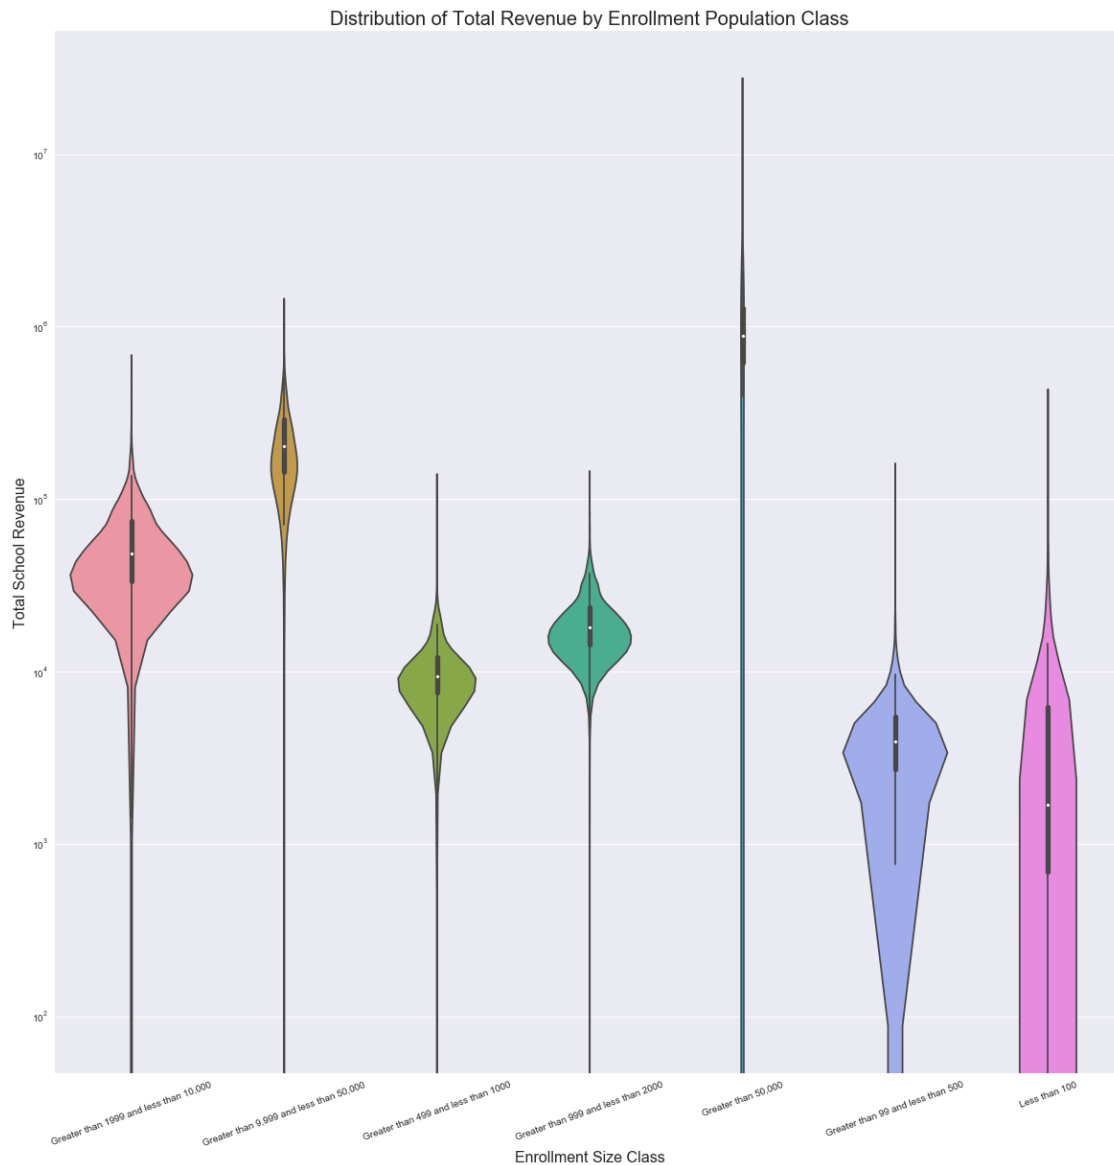
for i in range(0,7):
    plt.bar(i, df_pop_count["Enrollment Size Class"][i], label = list(df_pop_sum.index))
plt.legend()
plt.show()
plt.close()

```



It would be interesting to compare the enrollment population class to the distribution of total revenue for schools in that class, so that is shown in the following Violin Plot

```
In [18]: plt.figure(figsize=(20,20))
violinplot = sns.violinplot(x="Enrollment Size Class", y="TOTALREV", data=df, scale="cow")
violinplot.axes.set_yscale('log') #Make axis readable
plt.title("Distribution of Total Revenue by Enrollment Population Class", fontsize=20)
plt.xticks(rotation=20)
plt.ylabel('Total School Revenue', fontsize= 16)
plt.xlabel('Enrollment Size Class', fontsize= 16)
plt.show()
```



Another aspect of the dataset that could be explored is property tax revenue. It might be of interest to see how much property tax revenue schools in our dataset received. Here we see the frequency of schools within particular property tax revenue brackets.

```
In [19]: pt_mean = 2430 # Data obtained from https://www.urban.org/sites/default/files/publicati
pt_sd = 1599

property_taxes = df["T06"].dropna()
bin_names = ["Less than " + str(pt_mean - 2*pt_sd), str(pt_mean - 2*pt_sd) + " to " + s
pt_bins = pd.cut(property_taxes, [pt_mean-2*pt_sd, pt_mean-1*pt_sd, pt_mean, pt_mean+1*

pt_taxes = pd.DataFrame(property_taxes.groupby(pt_bins).count())
df["Property Tax Class"] = pt_bins
```

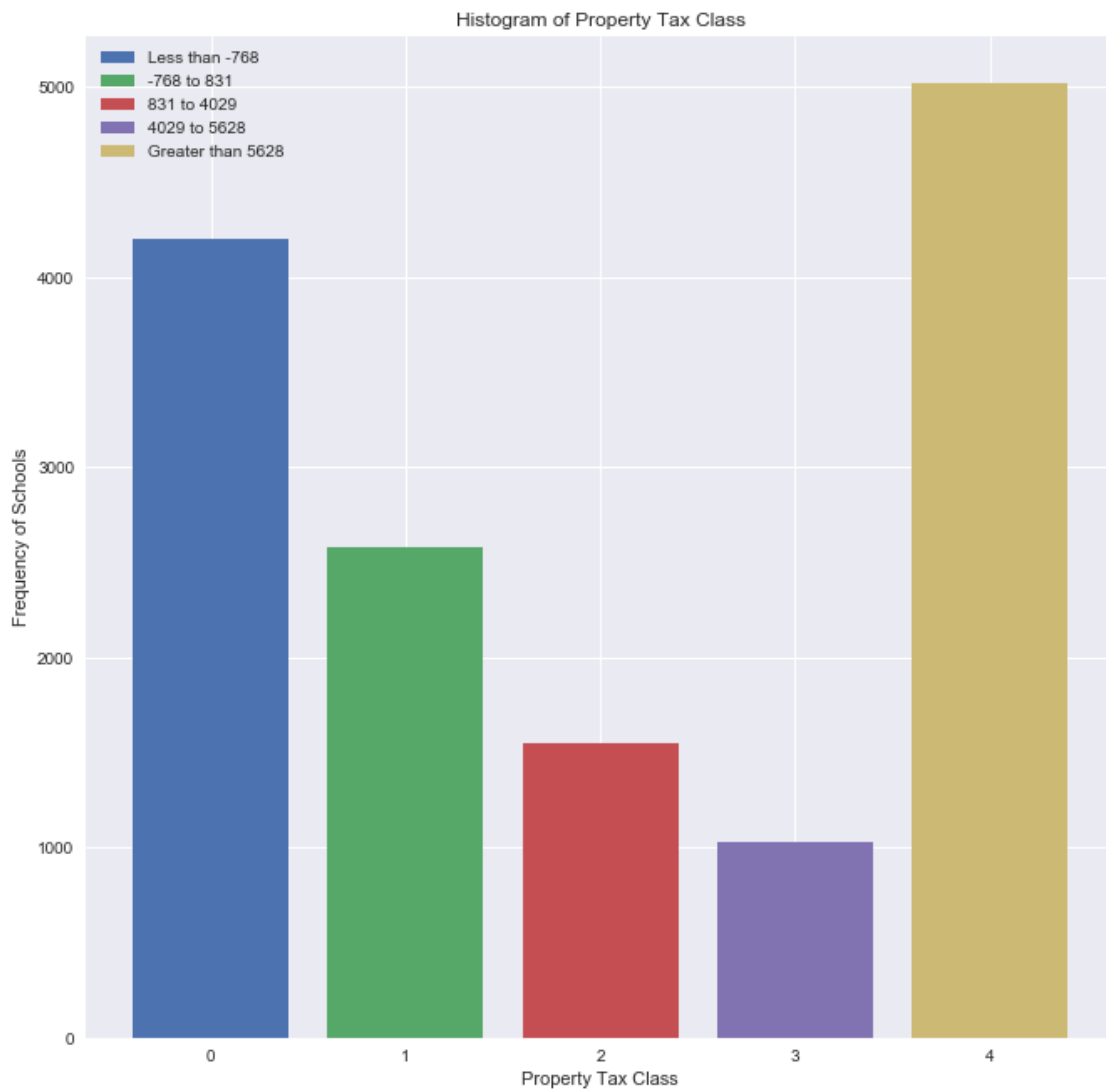
```

plt.figure(figsize=(11,11))

plt.xlabel('Property Tax Class')
plt.ylabel('Frequency of Schools')
plt.title('Histogram of Property Tax Class')

for i in range(0,5):
    plt.bar(i, pt_taxes["T06"][i], label = list(pt_taxes.index)[i])
plt.legend()
plt.show()
plt.close()

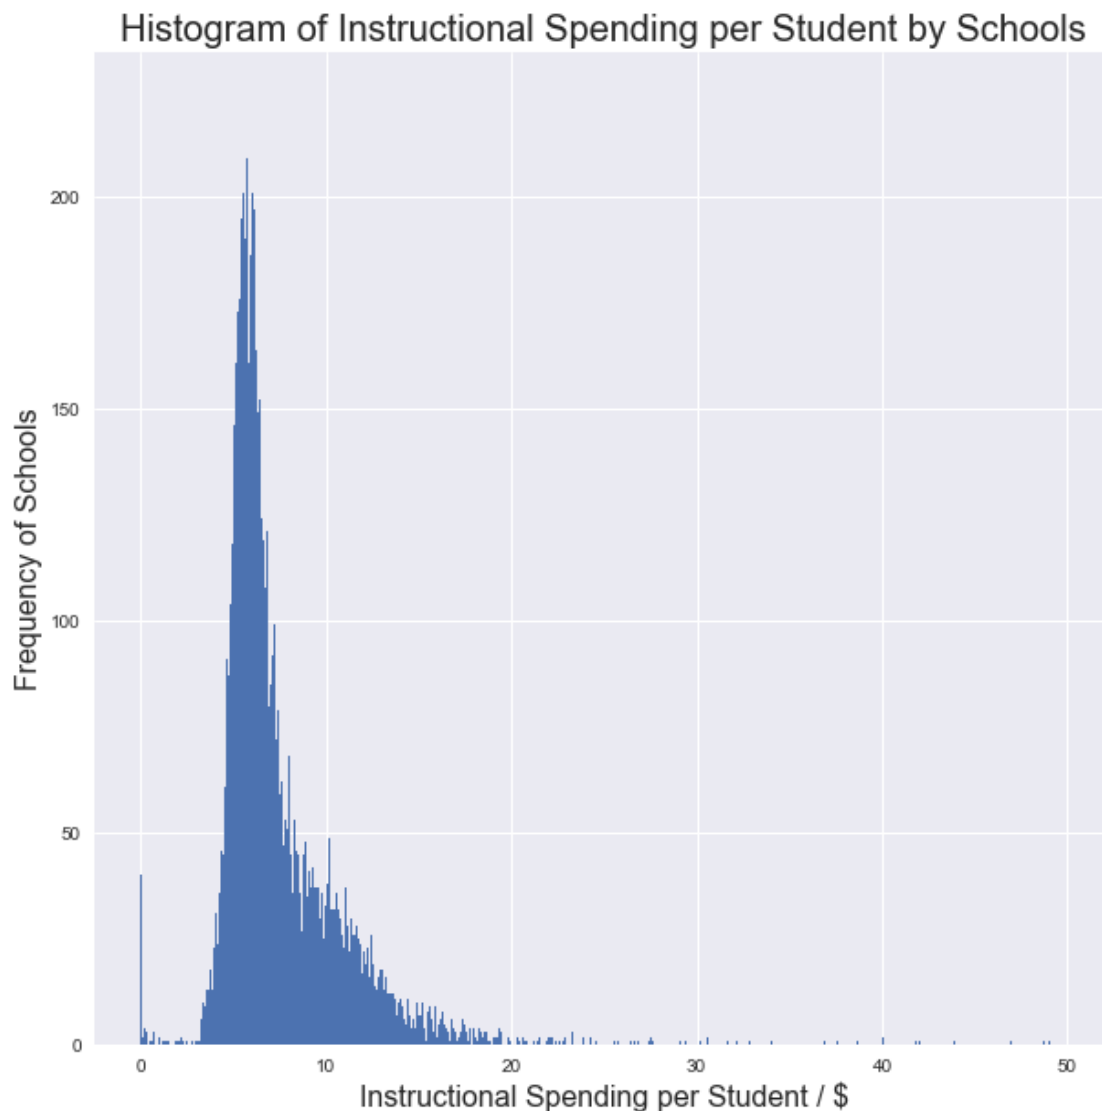
```



But perhaps the schools property tax revenue is not really indicative of very much the school could serve a lot of students or be in richer in poor communities. Perhaps a better aspect to look at

is is instructional spending per fall enrollment. Let's create a new column "Instructional spending per student" and get a histogram for that!

```
In [20]: df["Instructional spending per student"] = df["TCURINST"]/df[df["V33"] != 0]["V33"]
plt.figure(1, figsize=(10,10))
plt.xlabel('Instructional Spending per Student / $', fontsize = 16)
plt.ylabel('Frequency of Schools', fontsize = 16)
plt.title('Histogram of Instructional Spending per Student by Schools', fontsize = 20)
df["Instructional spending per student"].hist(range=[0, 50], bins = 1000)
plt.show()
plt.close()
```



## 2 Linear Regression!

That was fun! Now let's perform some regression. We'll perform regression using property tax to predict federal revenue for disadvantaged students, instructional spending, and school debt.

We'll build 3 models using simple linear regression

Model 1: Regressing on Property Taxes vs Federal Title 1 Funding (Disadvantaged Students)

Model 2: Regressing on Property Taxes vs Instructional Spending

Model 3: Regressing on Property Taxes vs Outstanding Debt at the start of the fiscal year

But before we train our model, we will remove outliers and zero value data in our property tax columns, as it's believed that they do not add the training of the model.

```
In [21]: # Remove outlier property taxes and 0 property taxes.
df_reg = df.copy()
df_reg = df_reg[(np.abs(stats.zscore(df_reg["T06"])) < 3) & (df_reg["T06"] != 0)]

taxes = df_reg["T06"]
title1 = df_reg["C14"]
spending = df_reg["TCURINST"]
debt = df_reg["_19H"]
```

Additionally, let us partition the data into training and validation sets, so we can see check how our models perform.

```
In [22]: # Partitioning validation and training data

taxes.shape
train_num = int(np.round(taxes.shape[0]*.8))

taxes_train = taxes[0:train_num]
taxes_valid = taxes[train_num:]
title1_train = title1[0:train_num]
title1_valid = title1[train_num:]
spending_train = spending[0:train_num]
spending_valid = spending[train_num:]
debt_train = debt[0:train_num]
debt_valid = debt[train_num:]
```

```
In [23]: plt.figure(1, figsize=(6,6))
myOLS_title1 = sm.OLS(title1,taxes).fit()
plt.plot(taxes, myOLS_title1.predict(taxes), color = 'red')
plt.scatter(taxes,title1)
plt.title("Property Taxes vs. Title 1 Funding", fontsize= 16)
plt.xlabel("Property Taxes", fontsize= 14)
plt.ylabel("Title 1 Funding", fontsize= 14)
plt.show()
plt.close()

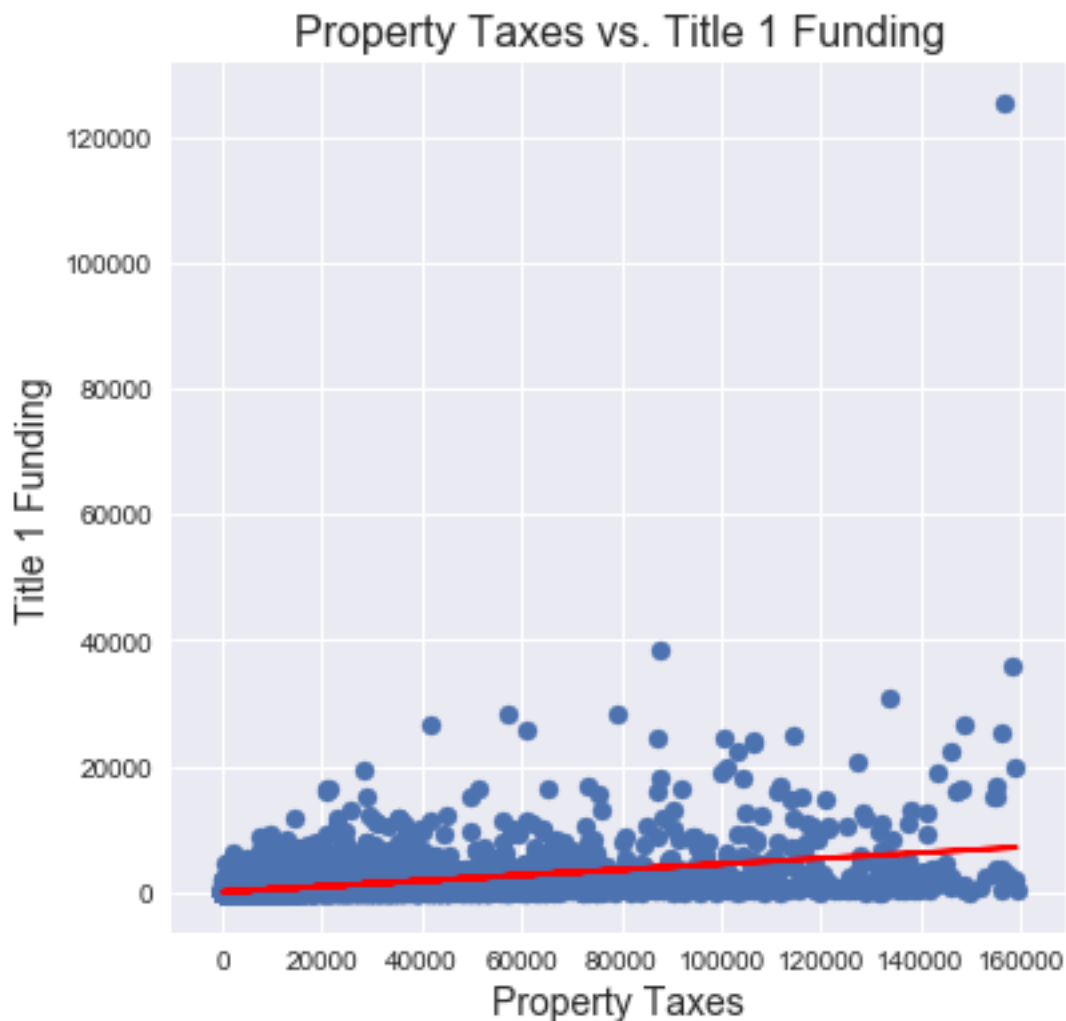
plt.figure(2, figsize=(6,6))
myOLS_spending = sm.OLS(spending,taxes).fit()
```

```

plt.plot(taxes, myOLS_spending.predict(taxes), color = 'red')
plt.scatter(taxes, spending)
plt.title("Property Taxes vs Instructional Spending", fontsize= 16)
plt.xlabel("Property Taxes", fontsize= 14)
plt.ylabel("Instructional Spending", fontsize= 14)
plt.show()
plt.close()

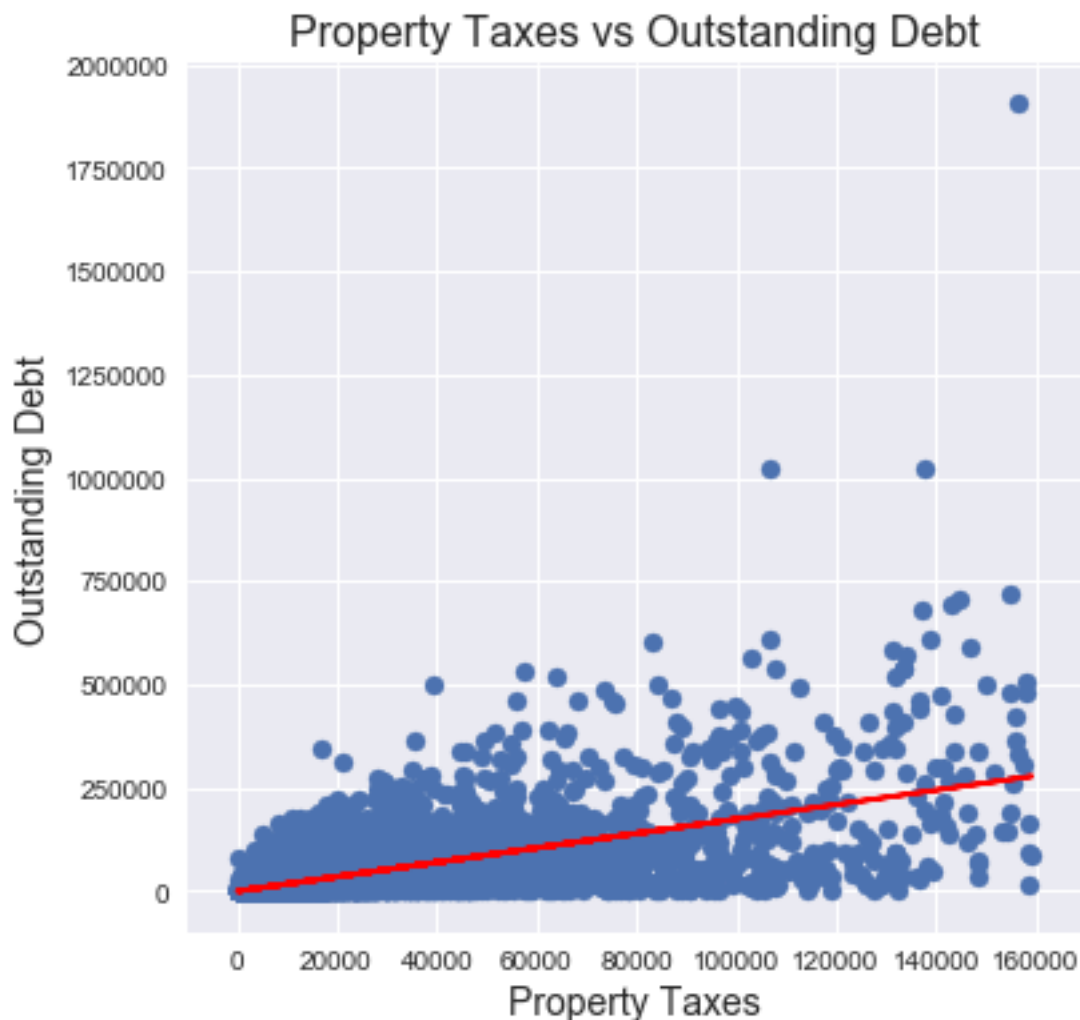
plt.figure(3, figsize = (6,6))
myOLS_debt = sm.OLS(debt, taxes).fit()
plt.plot(taxes, myOLS_debt.predict(taxes), color = 'red')
plt.scatter(taxes, debt)
plt.title("Property Taxes vs Outstanding Debt", fontsize= 16)
plt.xlabel("Property Taxes", fontsize= 14)
plt.ylabel("Outstanding Debt", fontsize= 14)
plt.show()
plt.close()

```









It looks as if property tax revenue best predicts instructional spending! Let's take a better look by calculating the mean squared errors.

```
In [24]: myOLS = sm.OLS(title1_train,taxes_train).fit()
         title1_hat = myOLS.predict(taxes_valid)
         title1_mse = 1/len(title1_valid)*np.dot((title1_valid - title1_hat),(title1_valid - title1_hat))
         print("The MSE for the model taxes~title1 is:", title1_mse)
         myOLS = sm.OLS(spending_train,taxes_train).fit()
         spending_hat = myOLS.predict(taxes_valid)
         spending_mse = 1/len(spending_valid)*np.dot((spending_valid - spending_hat),(spending_valid - spending_hat))
         print("The MSE for the model taxes~spending is:", spending_mse)
         myOLS = sm.OLS(debt_train,taxes_train).fit()
         debt_hat = myOLS.predict(taxes_valid)
         debt_mse = 1/len(debt_valid)*np.dot((debt_valid - debt_hat),(debt_valid - debt_hat))
         print("The MSE for the model taxes~debt is:", debt_mse)
         # Compare MSE
         print(min(title1_mse, spending_mse, debt_mse))
```

```
The MSE for the model taxes~title1 is: 2050414.15614
The MSE for the model taxes~spending is: 282908769.491
The MSE for the model taxes~debt is: 2616542250.57
2050414.15614
```

```
In [25]: myOLS = sm.OLS(title1_train,taxes_train).fit()
        title1_hat = myOLS.predict(taxes_valid)
        title1_mse = 1/len(title1_valid)*np.dot((title1_valid - title1_hat),(title1_valid - tit
        print("The MSE for the model taxes~title1 is: ", title1_mse)
        myOLS = sm.OLS(spending_train,taxes_train).fit()
        spending_hat = myOLS.predict(taxes_valid)
        spending_mse = 1/len(spending_valid)*np.dot((spending_valid - spending_hat),(spending_v
        print("The MSE for the model taxes~spending is:", spending_mse)
        myOLS = sm.OLS(debt_train,taxes_train).fit()
        debt_hat = myOLS.predict(taxes_valid)
        debt_mse = 1/len(debt_valid)*np.dot((debt_valid - debt_hat),(debt_valid - debt_hat))
        print("The MSE for the model taxes~debt is:      ", debt_mse)
```

```
The MSE for the model taxes~title1 is:    2050414.15614
The MSE for the model taxes~spending is:  282908769.491
The MSE for the model taxes~debt is:      2616542250.57
```

The mean squared error for title 1 and taxes is most minimum, however the best line fit looks best for debt and taxes

### 3 Logistic Regression!

Let's try logistic regression now! Perhaps it'll be interesting to see if property taxes, title 1 funding, funding for disabled students, funding for "math, science" teacher quality, and other forms of federal revenue-- are all predictors of whether a school is in debt or not.

Because logistic regression performs classification towards categorical data, let's create a categorical column debt/no debt.

```
In [26]: df["Debt Class"] = df["_19H"] == 0
        print("Percentage of schools with no debt: ", df[df["Debt Class"] == True]["_19H"].count)
        print("percentage of schools with debt: ", 1 - df[df["Debt Class"] == True]["_19H"].count)
```

```
Percentage of schools with no debt:  0.225792988314
percentage of schools with debt:    0.774207011686
```

```
In [27]: #Preprocess columns
        taxes = df["T06"]
        spending = df["TOTALEXP"]
        title1 = df["C14"]
        idea = df["C15"]
```

```

mstq = df["C16"]
fedrev = df["C20"]
debt = df["Debt Class"]

# Originally no normalization was performed, over flow occurred when exponentiating.
# Then a mean normalization was performed, resulting in many negative entries, which caused
# Finally, a min-max normalization was tried and worked.

scaler = MinMaxScaler()
categories = [taxes, spending, title1, idea, mstq, fedrev, debt]
scaler.fit(categories)
MinMaxScaler(copy=True, feature_range=(0, 1))

# We'll need to preprocess the columns we want to look at and partition the training and
data_x = np.column_stack(scaler.transform(categories))
data_y = np.reshape(debt, (data_x.shape[0], 1))

train_x, valid_x, train_y, valid_y = train_test_split(data_x, data_y, test_size=0.10, r

```

/usr/local/lib/python3.4/dist-packages/numpy/core/fromnumeric.py:57: FutureWarning:

reshape is deprecated and will raise in a subsequent release. Please use .values.reshape(...) in

We'll need to borrow the implementation of the logistic regression given to us in lecture

```

In [28]: def sigmoid(X, w):
        """
        Compute the elementwise sigmoid of the product Xw
        Data in X should be rows, weights are a column.
        """
        return 1 / (1 + np.exp(-np.dot(X, w)))

def gradient(X, y, w, onept, lamb=0):
    """
    Compute gradient of regularized loss function.
    Accomodate for if X is just one data point.
    """
    if onept:
        return 2 * lamb * w - ((y - sigmoid(X, w)) * X).reshape(w.size, 1)
    return 2 * lamb * w - np.dot(X.T, y - sigmoid(X, w)) / y.size

def loss(X, y, w, lamb=0):
    """
    Compute total loss for the data in X, labels in y, params w
    """

```

```

sumcost = 0
for i in range(X.shape[0]):
    sumcost += y[i] * np.log(sigmoid(X[i], w)) + (1 - y[i]) * np.log(1 - sigmoid(X[i], w))
return lamb * np.linalg.norm(w)**2 - sumcost / y.size

def accuracy(X, y, w):
    """
    Compute accuracy for data in X, labels in y, params w
    """
    results = np.round(sigmoid(X, w))
    score = sum([results[i] == y[i] for i in range(y.size)]) / y.size
    return score[0]

```

Training the model:

```

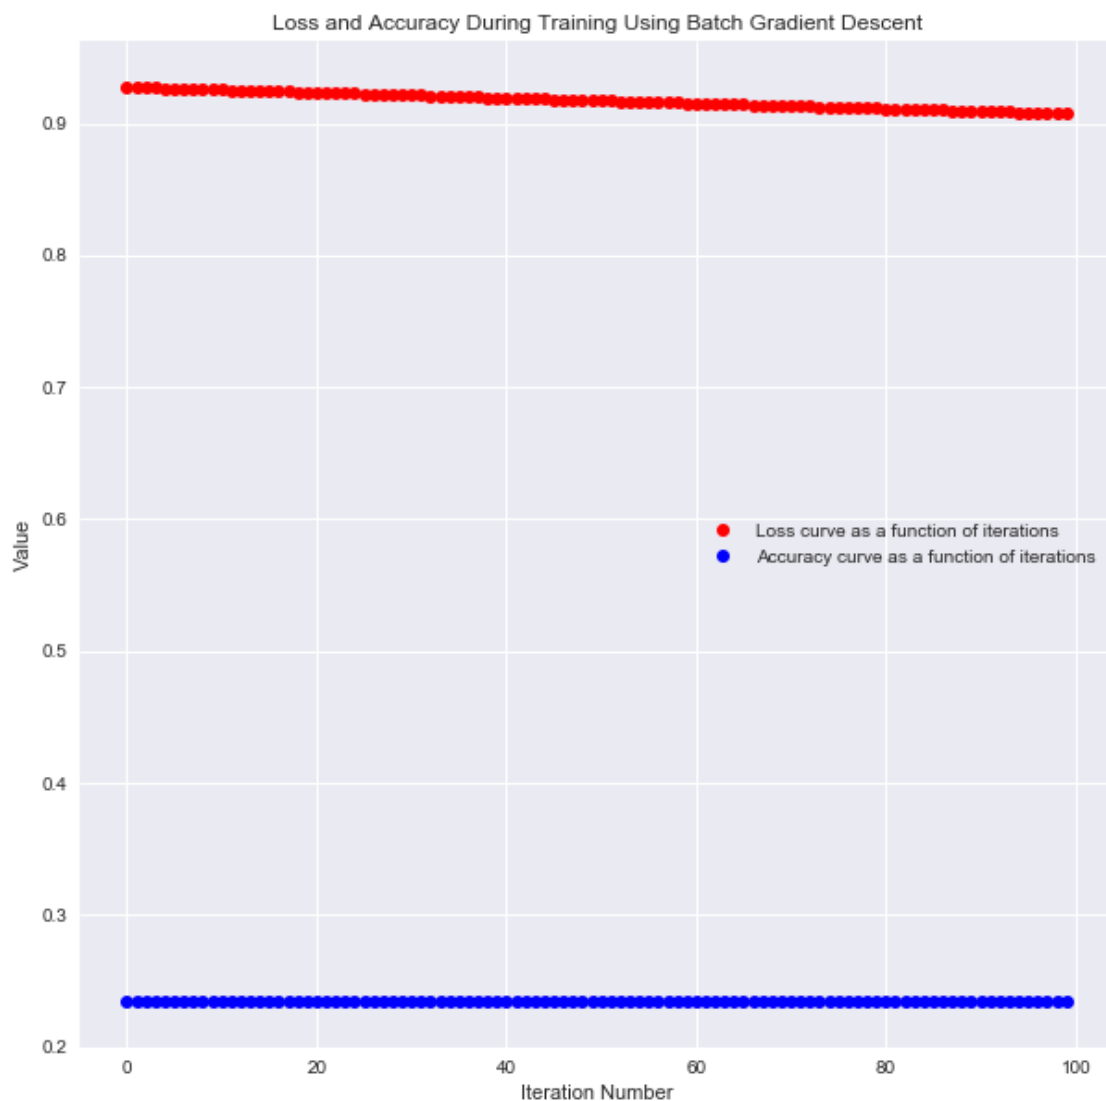
In [29]: weights = np.asarray([np.random.rand() for i in range(train_x.shape[1])]).reshape(train_x.shape[1])
weights /= np.linalg.norm(weights)
losses = []
accuracies = []
epsilon = 0.001
num_iterations = 100

for i in range(100):
    diff = epsilon * gradient(train_x, train_y, weights, False)
    weights = weights - diff
    losses.append(loss(train_x, train_y, weights))
    accuracies.append(accuracy(valid_x, valid_y, weights))

In [30]: plt.figure(figsize=[10,10])
loss_line = plt.plot(np.arange(num_iterations), losses, 'ro', label="Loss curve as a function of iteration number")
accuracy_line = plt.plot(np.arange(num_iterations), accuracies, 'bo', label = "Accuracy curve as a function of iteration number")
plt.legend()
plt.title('Loss and Accuracy During Training Using Batch Gradient Descent')
plt.ylabel('Value')
plt.xlabel('Iteration Number')

plt.show()
plt.close()
print(accuracy(valid_x, valid_y, weights))

```



0.234353268428

## 4 Concluding visualization

As a fun conclusion, we made a corruption index, which is defined for a set of data as:

$$\frac{\sum_{i=0}^{\#rows} \text{"Salaries and employee benefits for general and school administration"}}{\sum_{i=0}^{\#rows} \text{"Total Revenue"}}$$

This is defined as corruption as this money is going into to the pockets of the school administrators, and not the school teachers or students.

```
In [31]: # GENERATES OUR CORRUPTION INDEX PER STATE CHLOROPETH
```

```
import plotly.plotly as py
from plotly.graph_objs import *

scl = [[0.0, 'rgb(0,255,0)'],[1.0, 'rgb(188,0,0)']]
data = [ dict(
    type='choropleth',
    colorscale = scl,
    autocolorscale = False,
    locations = States['State'],
    z = States['Corruption'].astype(float),
    locationmode = 'USA-states',
    marker = dict(
        line = dict (
            color = 'rgb(255,255,255)',
            width = 2
        ) ),
    colorbar = dict(
        title = "Corruption Index")
    ) ]
layout = Layout(
    autosize=False,
    geo=dict(
        countrycolor='rgb(102, 102, 102)',
        countrywidth=0.1,
        lakecolor='rgb(255, 255, 255)',
        landcolor='rgba(237, 247, 138, 0.28)',
        lonaxis=dict(
            gridwidth=1.5999999999999999,
            range=[-180, -50],
            showgrid=False
        ),
        projection=dict(
            type='albers usa'
        ),
        scope='usa',
        showland=True,
        showrivers=False,
        showsubunits=True,
        subunitcolor='rgb(102, 102, 102)',
        subunitwidth=0.5
    ),
    hovermode='closest',
    images=list([
        dict(
            x=1,
            y=0.6,
            sizex=0.155,
```

```

        sizey=0.4,
        source='http://i.imgur.com/Xe3f1zg.png',
        xanchor='right',
        xref='paper',
        yanchor='bottom',
        yref='paper'
    )
]),
showlegend=True,
title='<b>Corruption Index per state</b>',
width= 800,
margin = dict(
    l=0,
    r=50,
    b=100,
    t=100,
    pad=4)
)
fig = Figure(data=data, layout=layout)
py.ipplot(fig, filename='DSD_2')

```

Out[31]: <plotly.tools.PlotlyDisplay object>

It appears that Mississippi and Oklahoma have the most corruption in their school boards. 10% of their revenue is going towards the school administrators. *Please note that this metric is just for fun and should not be taken seriously. There are issues in measuring corruption this way, since schools need administrators in order to run.*